


# Loss functions



# Loss functions

*How good is your neural network?*

# Loss functions

*How good is your neural network?*

## Inputs


e.g.

- Image (pixel values)
- Sentence (encoded)
- State (for RL)

## Targets

e.g.

- 'Cat'/'Dog' (encoded)
- Next word (encoded)
- TD(0) (for RL)



# Loss functions

## *Some pedantry*

***Loss function** and **Cost function** are often used interchangeably but they are different:*

- The **loss function** measures the network performance **on a single datapoint***
- The **cost function** is the **average of the losses** over the entire training dataset*

*Our goal is to **minimize the cost function**.*

*In reality, we actually use **batch losses** as a proxy for the cost function.*

# Loss/Cost functions


## *Some stuff to remember*

- The cost function distils the performance of the network down to a **single scalar number**.
- Generally (although not necessarily) **loss**  $\geq 0$  (so **cost**  $\geq 0$  also)
  - $\text{cost}=0 \Rightarrow$  perfect performance (on the given dataset)
- The **lower the value of the cost function, the better the performance** of the network.
  - Hence gradient descent
- A cost function must **faithfully represent the purpose** of the network.

# Loss/Cost functions

## *Some stuff to remember*

- The cost function distils the performance of the network down to a **single scalar number**.
- Generally (although not necessarily) **loss**  $\geq 0$  (so **cost**  $\geq 0$  also)
  - $\text{cost}=0 \Rightarrow$  perfect performance (on the given dataset)
- The **lower the value of the cost function, the better the performance** of the network.
  - Hence gradient descent
- A cost function must **faithfully represent the purpose** of the network.



# Loss/Cost functions

## *Loss functions*

- **Classification**
  - *Maximum likelihood*
  - *Binary cross-entropy (aka log loss)*
  - *Categorical cross-entropy*
- **Regression** (i.e. function approximation)
  - *Mean Squared Error*
  - *Mean Absolute Error*
  - *Huber Loss*

# Classification loss functions



## *Loss functions*

- **Classification**
  - *Maximum likelihood*
  - *Binary cross-entropy (aka log loss)*
  - *Categorical cross-entropy*



# Classification loss functions

## *Maximum likelihood*





Target

(1, 1, 0, 0, 0)

Cat -  
Dog -  
Bear -  
Hitler -  
Putin -

# Classification loss functions

## *Maximum likelihood*



$$Loss = -\frac{1}{N} \sum_{i=1}^N y_i \cdot (p(y_i)) + (1 - y_i) \cdot (1 - p(y_i))$$

$$\begin{aligned} &= -1/5 [ (1*0.5 + 0*0.5) + (1*0.9 + 0*0.1) + (0*0.2 + 1*0.8) \\ &\quad + (0*0.6 + 1*0.4) + (0*0.1 + 1*0.9) ] \\ &= -0.7 \end{aligned}$$


**Target**

(1, 1, 0, 0, 0)  
Cat -  
Dog -  
Bear -  
Hitler -  
Putin -

# Classification loss functions

*Binary cross-entropy (aka Log loss)*

Torch.nn.BCELoss



$$Loss = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$


**Target**

(0,	0,	1,	0,	1)
Cat	Dog	Bear	Hitler	Putin
-	-	-	-	-

# Classification loss functions

*Binary cross-entropy (aka Log loss)*

Torch.nn.BCELoss



$$Loss = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

Target



(0, 0, 1, 0, 1)  
Cat -  
Dog -  
Bear -  
Hitler -  
Putin -

$$\begin{aligned} &= -1/5 [ (1 * \log(0.8)) + (1 * \log(0.6)) + (1 * \log(0.8)) \\ &\quad + (1 * \log(0.7)) + (1 * \log(0.9)) ] \\ &= 0.28 \end{aligned}$$

# Classification loss functions

*Binary cross-entropy (aka Log loss)*

Torch.nn.BCELoss



$$Loss = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

**Target**

(0, 0, 1, 0, 1)  
Cat -  
Dog -  
Bear -  
Hitler -  
Putin -

$$= -1/5 [ (1 * \log(0.8)) + (1 * \log(0.6)) + (1 * \log(0.8)) + (1 * \log(0.7)) + (1 * \log(0.9)) ]$$

$$= 0.28$$

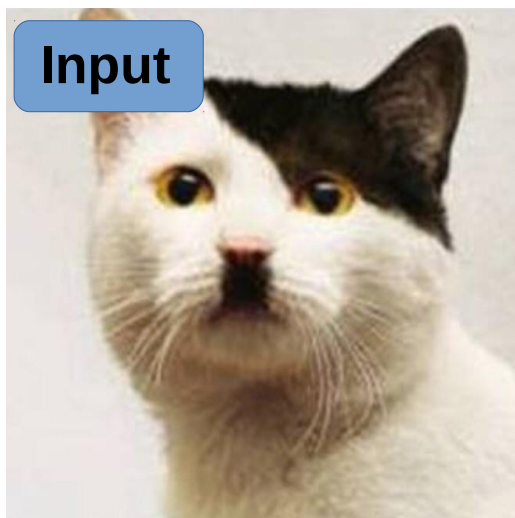


# Classification loss functions


## Categorical cross-entropy

Torch.nn.NLLLoss

Torch.nn.CrossEntropyLoss



Input



Cat Dog Bear Hitler Putin  
(.8, 0, 0, .2, 0)

This one does the softmax for you

$$= \log \left( \frac{\exp(x_i)}{\sum_j \exp(x_j)} \right)$$

Must be a probability distribution

$$Loss = - \sum_{i=1}^N y_i \cdot \log(p(y_i))$$

$$= - [ 1 * \log(0.8) ]$$

$$= 0.22$$

Target

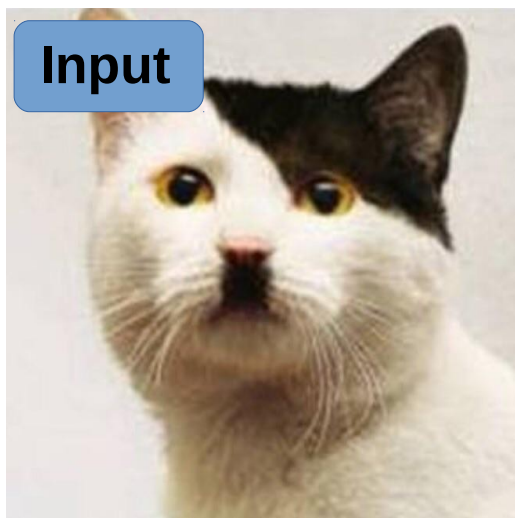
(1, 0, 0, 0, 0)  
Cat -  
Dog -  
Bear -  
Hitler -  
Putin -

# Classification loss functions


## Categorical cross-entropy

Torch.nn.NLLLoss

Torch.nn.CrossEntropyLoss



Input



Cat - Dog - Bear - Hitler - Putin  
(.8, 0, 0, .2, 0)

This one does the softmax for you

$$= \log \left( \frac{\exp(x_i)}{\sum_j \exp(x_j)} \right)$$

Must be a probability distribution

$$Loss = - \sum_{i=1}^N y_i \cdot \log(p(y_i))$$

$$= - [ 1 * \log(0.8) ]$$

$$= 0.22$$

Target

(1, 0, 0, 0, 0)  
Cat - Dog - Bear - Hitler - Putin

This loss function can also be used when you have a single binary (e.g. yes/no) output

# Regression loss functions

## *Loss functions*

- **Regression** (i.e. function approximation)
  - *Mean Squared Error*
  - *Mean Absolute Error*
  - *Huber Loss*




# Regression loss functions

aka L2 loss

*Mean squared error*

Torch.nn.MSELoss

Input  $x$



Target  $y$


$$Loss = (y - f(x))^2$$

# Regression loss functions


aka L2 loss

*Mean squared error*

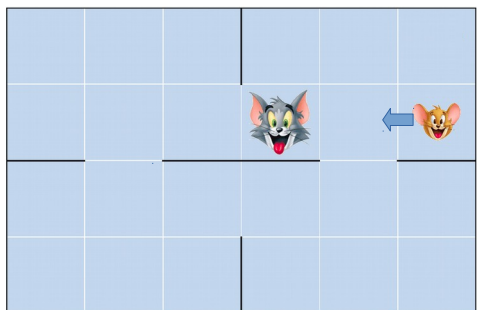
Torch.nn.MSELoss



state  $S_t$  & action  $A_t$




reward  $R_{t+1}$



state  $S_{t+1}$  & [best] action  $A_{t+1}$

**Target**

$$R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$$



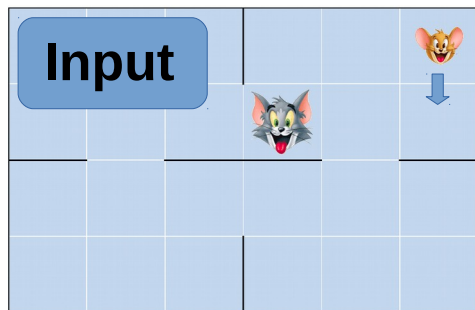
# Regression loss functions

aka L2 loss


*Mean squared error*

Torch.nn.MSELoss

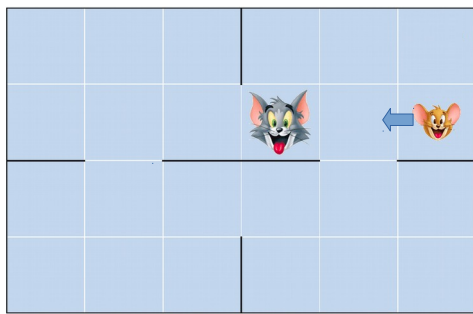
$$\begin{aligned} \text{Loss} &= (y - f(x))^2 \\ &= \left( (R_{t+1} + \gamma \max_a Q(S_{t+1}, a)) - Q(S_t, A_t) \right)^2 \end{aligned}$$



state  $S_t$  & action  $A_t$




reward  $R_{t+1}$



state  $S_{t+1}$  & [best] action  $A_{t+1}$

**Target**

$$R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$$




# Regression loss functions

aka L1 loss

*Absolute error*

Torch.nn.L1Loss

Input  $x$



Target  $y$


$$Loss = |y - f(x)|$$

# Regression loss

*Huber loss*

Torch.nn.SmoothL1Loss

Input  $x$



Target  $y$

$$Loss = \begin{cases} \frac{1}{2}(y - f(x))^2, & \text{if } |y - f(x)| \leq \delta \\ \delta|y - f(x)| - \frac{1}{2}\delta^2, & \text{otherwise} \end{cases}$$

Huber loss combines the best features of MSE and AE:

- like MSE for small errors.
- like AE otherwise.

This avoids over-shooting:

- MSE has a **very steep gradient** when the error is **large** (i.e. for outliers) because of its quadratic form.
- AE has a **steep gradient** when the error is **small** because it doesn't have a 'flat' bottom.

