


Recurrent Neural Networks (RNN)

A family of neural architectures


Core idea: Apply the same weights W repeatedly



A Recurrent Neural Network (RNN). Three time-steps are shown.

Unlike the conventional translation models, where only a finite window of previous words would be considered for conditioning the language model, **Recurrent Neural Networks (RNN)** are capable of conditioning the model on all previous words in the corpus.

RNN architecture where each vertical rectangular box is a hidden layer at a time-step, t . Each such layer holds a number of neurons, each of which performs a linear matrix operation on its inputs followed by a non-linear operation (e.g. $\tanh()$). At each time-step, there are two inputs to the hidden layer: the output of the previous layer h_{t-1} , and the input at that timestep x_t . The former input is multiplied by a weight matrix $W^{(hh)}$ and the latter by a weight matrix $W^{(hx)}$ to produce output features h_t , which are multiplied with a weight matrix $W^{(S)}$ and run through a softmax over the vocabulary to obtain a prediction output \hat{y} of the next word




The inputs and outputs to a neuron of a RNN

$$h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_{[t]})$$

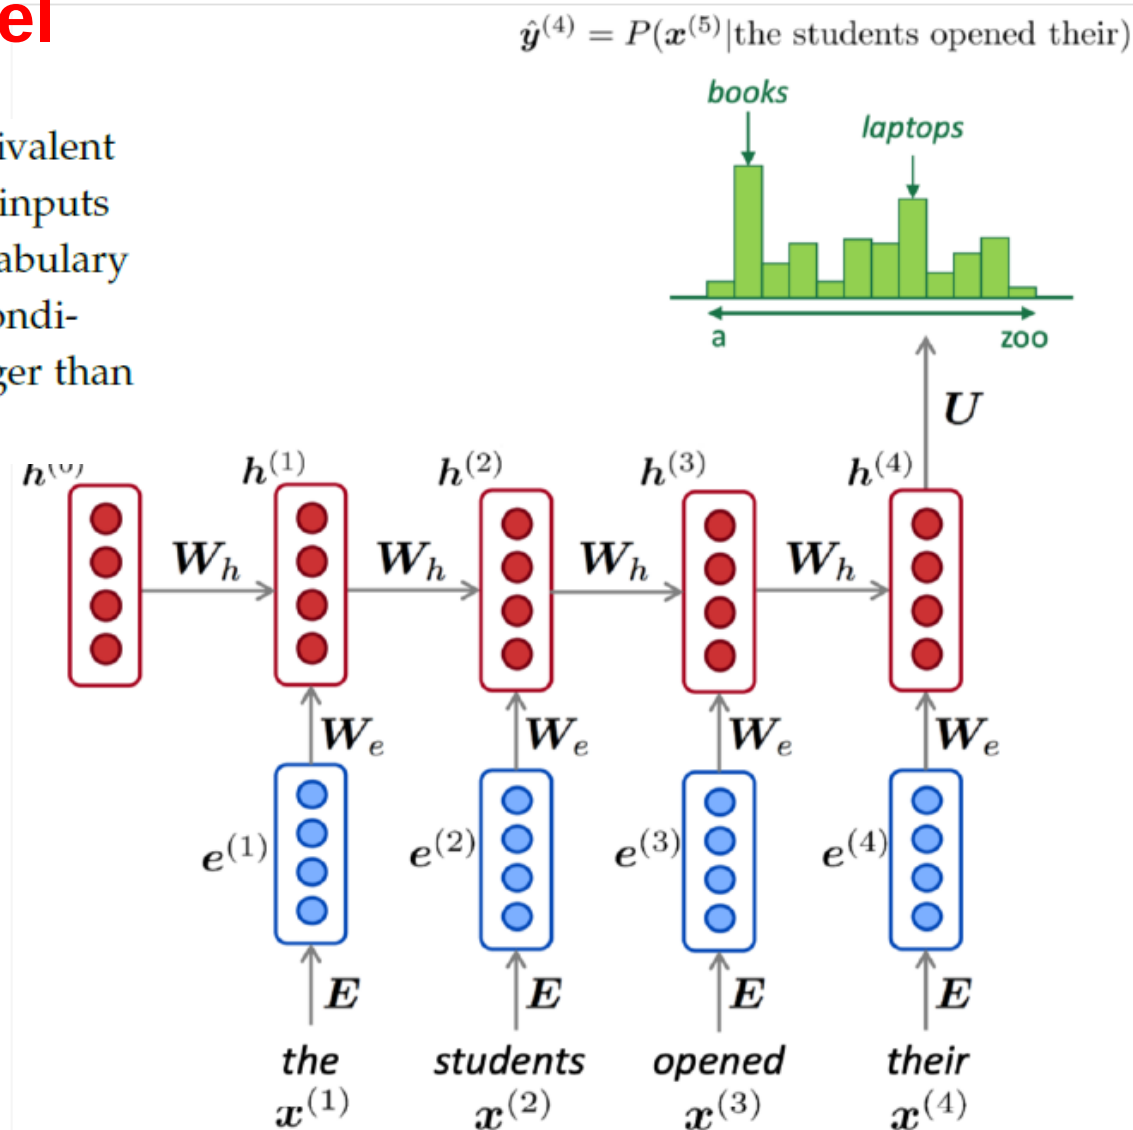
$$\hat{y}_t = \text{softmax}(W^{(S)}h_t)$$

What is interesting here is that the same weights $W^{(hh)}$ and $W^{(hx)}$ are applied repeatedly at each timestep. Thus, the number of parameters the model has to learn is less, and most importantly, is independent of the length of the input sequence - thus defeating the curse of dimensionality!




An RNN Language Model

The notation in this image is slightly different: here, the equivalent of $W^{(hh)}$ is W_h , $W^{(hx)}$ is W_e , and $W^{(S)}$ is U . E converts word inputs $x^{(t)}$ to word embeddings $e^{(t)}$. The final softmax over the vocabulary shows us the probability of various options for token $x^{(5)}$, conditioned on all previous tokens. The input could be much longer than 4-5 tokens.




The Recurrent Neuron

- x_t : Input at time t
- h_{t-1} : State at time $t-1$




$$h_t = f(W_h h_{t-1} + W_x x_t)$$

Simple recurrent neural network illustrated as a feedforward network.



A simple recurrent neural network shown unrolled in time.

Network layers are copied for each time step, while the weights U , V and W are shared in common across all time steps.



A Simple RNN Language Model

output distribution

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{U}\mathbf{h}^{(t)} + \mathbf{b}_2) \in \mathbb{R}^{|V|}$$

hidden states

$$\mathbf{h}^{(t)} = \sigma(\mathbf{W}_h \mathbf{h}^{(t-1)} + \mathbf{W}_e \mathbf{e}^{(t)} + \mathbf{b}_1)$$

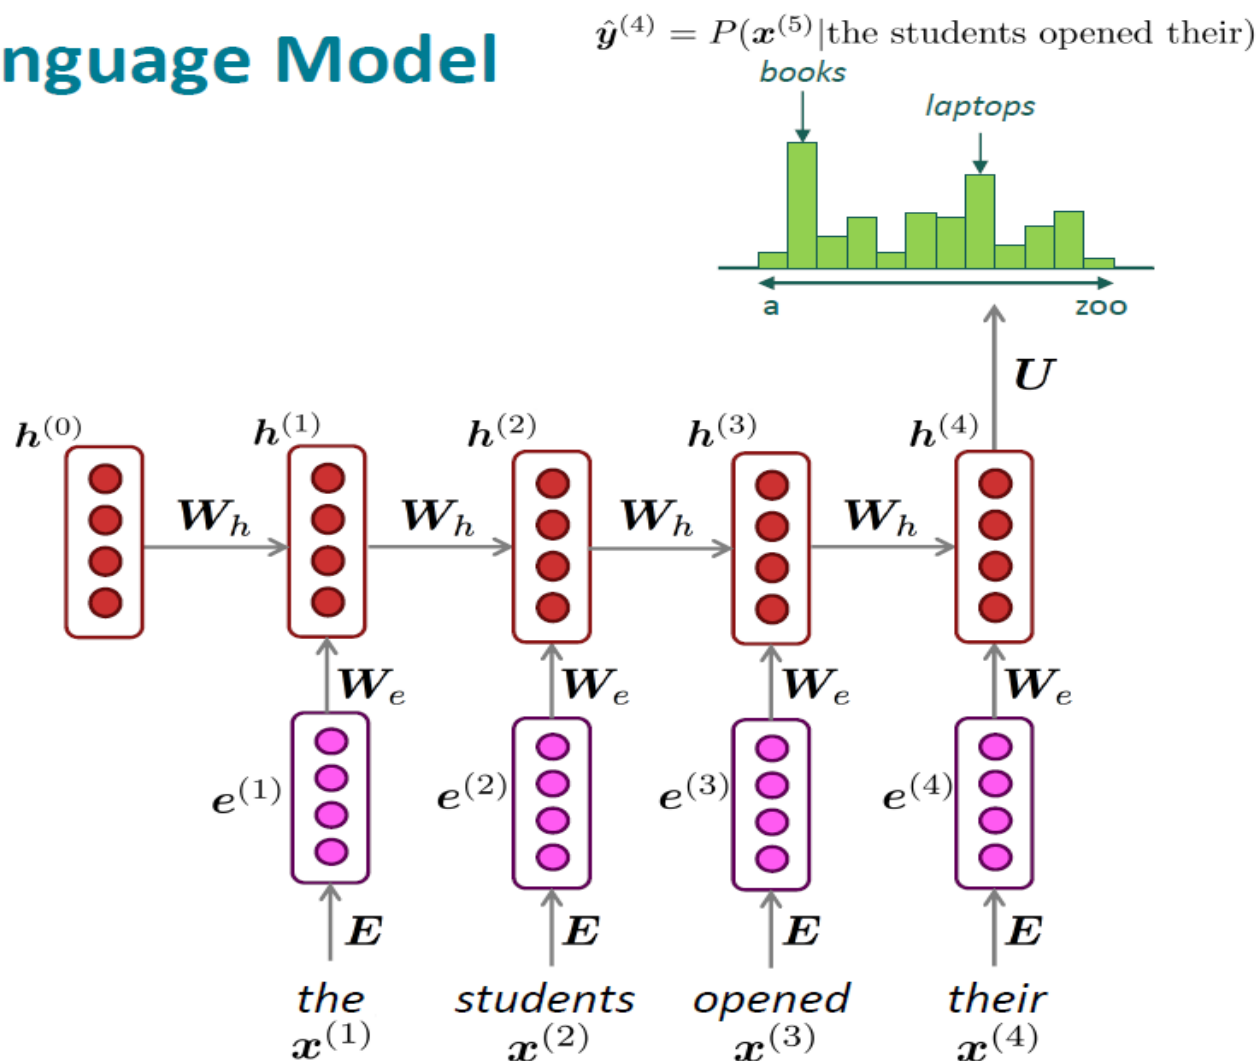
$\mathbf{h}^{(0)}$ is the initial hidden state

word embeddings

$$\mathbf{e}^{(t)} = \mathbf{E}\mathbf{x}^{(t)}$$

words / one-hot vectors

$$\mathbf{x}^{(t)} \in \mathbb{R}^{|V|}$$



Note: this input sequence could be much longer now!


RNN Language Models

RNN Advantages:

- Can process **any length** input
- Computation for step t can (in theory) use information from **many steps back**
- **Model size doesn't increase** for longer input context
- Same weights applied on every timestep, so there is **symmetry** in how inputs are processed.

RNN Disadvantages:

- Recurrent computation is **slow**
- In practice, difficult to access information from **many steps back**



- **The amount of memory required to run a layer of RNN is proportional to the number of words in the corpus.**
- We can consider a sentence as a minibatch, and a sentence with k words would have k word vectors to be stored in memory.
- Also, the RNN must maintain two pairs of W, b matrices. As aforementioned, while the size of W could be very large, it does not scale with the size of the corpus (unlike the traditional language models).
- For a RNN with 1000 recurrent layers, the matrix would be 1000×1000 regardless of the corpus size.